

GEOFF PLEISS, CURRICULUM VITAE

Postdoctoral Research Scientist
Department of Statistics and Zuckerman Institute
Columbia University
PI: Dr. John P. Cunningham

Phone: (408) 338-8247
Email: gmp2162@columbia.edu
Homepage: <http://www.geoffpleiss.com>
(last updated February 16, 2022)

EDUCATION

Ph.D. Computer Science (2020)
Cornell University (Ithaca, NY)
Committee: Kilian Q. Weinberger, Andrew Gordon Wilson, Karthik Sridharan
Dissertation: A Scalable and Flexible Framework for Gaussian Processes via Matrix-Vector Multiplication

M.Sc. Computer Science (2018)
Cornell University (Ithaca, NY)

B.Sc. Engineering: Computing with Applied Mathematics (2013)
Olin College of Engineering (Needham, MA)

EMPLOYMENT

2020– Postdoctoral Research Scientist
Columbia University (New York, NY)

2019-2020 Research Intern
ASAPP Inc. (Ithaca, NY)

2018 Research Intern
Microsoft Inc. (Redmond, WA)

2013–2015 Software Engineer
Pivotal Inc. (New York, NY)

CITATION STATISTICS

All statistics are based on Google Scholar, with manual corrections for errors.
Total citations of all publications: 4600+
Total citations of top-three most cited publications: 3200+
Publications (including technical reports) with 100+ citations: 9
Publications (including technical reports) with 10+ citations: 16

HONORS AND AWARDS

2022 AISTATS Top Reviewer (top 10%)
2020 NeurIPS Top Reviewer (top 10%)
2017 National Science Foundation Graduate Research Fellowship (honorable mention)
2016 National Science Foundation Graduate Research Fellowship (honorable mention)
2012 Barry M. Goldwater Scholarship (honorable mention)
2009–2013 Olin Merit Scholarship (full-tuition recipient)

INVITED TALKS

Oct. 2021	Dagstuhl Seminar on Probabilistic Numerics
Oct. 2021	University College London
Dec. 2020	Machine Learning for Nuclear Data Workshop
Sept. 2019	Bill and Melinda Gates Foundation
May 2019	Uber AI Symposium on Bayesian Optimization
Dec. 2018	Neural Information Processing Systems (Spotlight)
Jul. 2018	International Conference on Machine Learning
Aug. 2017	International Conference on Machine Learning

TEACHING ASSISTANTSHIPS

Fall 2017	CS6780 — <i>Advanced Topics in Machine Learning</i> Cornell University (w/ Prof. Kilian Weinberger)
Fall 2016	CS4786 — <i>Machine Learning for Data Science</i> Cornell University (w/ Prof. Karthik Sridharan)
Fall 2015	CS4700 — <i>Foundations of Artificial Intelligence</i> Cornell University (w/ Prof. Bart Selman)
Spring 2013	SCI3130 — <i>Advanced Classical Mechanics</i> Olin College (w/ Prof. Yevgeniya V. Zastavker)
Spring 2011	SCI1121 — <i>Computational Electricity and Magnetism</i> Olin College (w/ Prof. Mark Somerville)

PROFESSIONAL SERVICE

Area Chair	ICML (2022)
Reviewing	JMLR (2019–2021)
Reviewing	TPAMI (2020–2021)
Reviewing	AAAI (2017)
Reviewing	AISTATS (2019–2022)
Reviewing	ICLR (2022)
Reviewing	ICML (2019–2021)
Reviewing	NeurIPS (2018–2021)
Reviewing	UAI (2018)

MENTORSHIP AND OUTREACH

Fall 2020	LatinX in AI NeurIPS mentorship program
Spring 2018	Cornell “Expand Your Horizons” (STEM workshop for middle school girls)
Spring 2017	Cornell “GRASSHOPR” (After-school CS class at local middle school)
Spring 2016	Cornell “Expand Your Horizons”
Spring 2016	“Code4Kids” (After-school CS class at local elementary school)

PUBLICATIONS

* denotes equal author contribution (shared first-authorship).

Refereed Conference Publications

In machine learning, conferences are considered prestigious venues for publication. All venues listed here are highly selective (acceptance rate $\leq 25\%$) and have peer-reviewing and refereeing processes similar to journals.

- [C1] **Geoff Pleiss** and John P. Cunningham. The limitations of large width in neural networks: A deep Gaussian process perspective. In *Neural Information Processing Systems*, 2021.
- [C2] Anthony L. Caterini, Gabriel Loaiza-Ganem, **Geoff Pleiss**, and John P. Cunningham. Rectangular flows for manifold learning. In *Neural Information Processing Systems*, 2021.
- [C3] Andres Potapczynski, Luhuan Wu, Dan Biderman, **Geoff Pleiss**, and John P. Cunningham. Bias-free scalable Gaussian processes via randomized truncations. In *International Conference on Machine Learning*, 2021.
- [C4] Luhuan Wu, Andrew Miller, Lauren Anderson, **Geoff Pleiss**, David Blei, and John P. Cunningham. Hierarchical inducing point Gaussian process for inter-domain observations. In *Artificial Intelligence and Statistics*, 2021.
- [C5] **Geoff Pleiss**, Martin Jankowiak, David Eriksson, Anil Damle, and Jacob R. Gardner. Fast matrix square roots with applications to Gaussian processes and Bayesian optimization. In *Neural Information Processing Systems*, 2020.
- [C6] **Geoff Pleiss**, Tianyi Zhang, Ethan Elenberg, and Kilian Q. Weinberger. Identifying mislabeled data using the area under the margin ranking. In *Neural Information Processing Systems*, 2020.
- [C7] Martin Jankowiak, **Geoff Pleiss**, and Jacob R. Gardner. Deep sigma point processes. In *Uncertainty in Artificial Intelligence*, 2020.
- [C8] Martin Jankowiak, **Geoff Pleiss**, and Jacob R. Gardner. Parametric Gaussian process regressors. In *International Conference on Machine Learning*, 2020.
- [C9] Yurong You, Yan Wang, Wei-Lun Chao, Divyansh Garg, **Geoff Pleiss**, Bharath Hariharan, Mark Campbell, and Kilian Q. Weinberger. Pseudo-lidar++: Accurate depth for 3d object detection in autonomous driving. In *International Conference on Learning Representations*, 2020.
- [C10] Ke Wang*, **Geoff Pleiss***, Jacob R. Gardner, Stephen Tyree, Kilian Q. Weinberger, and Andrew Gordon Wilson. Exact Gaussian processes on a million data points. In *Neural Information Processing Systems*, 2019.
- [C11] Jacob R. Gardner*, **Geoff Pleiss***, David Bindel, Kilian Q. Weinberger, and Andrew Gordon Wilson. GPyTorch: Blackbox matrix-matrix Gaussian process inference with GPU acceleration. In *Neural Information Processing Systems*, 2018.
- [C12] **Geoff Pleiss**, Jacob R. Gardner, Andrew Gordon Wilson, and Kilian Q. Weinberger. Constant time predictive distributions for Gaussian processes. In *International Conference on Machine Learning*, 2018.
- [C13] Jacob R. Gardner, **Geoff Pleiss**, Ruihan Wu, Andrew Gordon Wilson, and Kilian Q. Weinberger. Product kernel interpolation for scalable Gaussian processes. In *Artificial Intelligence and Statistics*, 2018.
- [C14] **Geoff Pleiss***, Manish Raghavan*, Felix Wu, Jon Kleinberg, and Kilian Q. Weinberger. On fairness and calibration. In *Neural Information Processing Systems*, 2017.

- [C15] Chuan Guo*, **Geoff Pleiss***, Yu Sun*, and Kilian Q. Weinberg. On calibration of modern neural networks. In *International Conference on Machine Learning*, 2017.
- [C16] Paul Upchurch, Jacob Gardner, **Geoff Pleiss**, Kavita Bala, Robert Pless, Noah Snavely, and Kilian Q. Weinberger. Deep feature interpolation for image content changes. In *Computer Vision and Pattern Recognition*, 2017.
- [C17] Gao Huang, Yixuan Li, **Geoff Pleiss**, Zhuang Liu, John E. Hopcroft, and Kilian Q. Weinberger. Snapshot ensembles: Train 1, get M for free. In *International Conference on Learning Representations*, 2017.

Refereed Journal Publications

- [J1] Gao Huang, Zhuang Liu, **Geoff Pleiss**, Laurens van der Maaten, and Kilian Q. Weinberger. Convolutional networks with dense connectivity. *Pattern Analysis and Machine Intelligence*, 2019.
- [J2] James Knighton, **Geoff Pleiss**, Elizabeth Carter, Steven Lyon, M. Todd Walter, and Scott Steinschneider. Potential predictability of regional precipitation and discharge extremes using synoptic-scale climate information via machine learning: An evaluation for the eastern continental united states. *Journal of Hydrometeorology*, 20(5):883–900, 2019.

Preprints Under Submission

- [U1] Taiga Abe, E. Kelly Buchanan, **Geoff Pleiss**, Richard Zemel, and John P. Cunningham. Deep ensembles work, but are they necessary? *arXiv preprint arXiv:2202.06985*, 2022.
- [U2] Luhuan Wu, **Geoff Pleiss**, and John P. Cunningham. Variational nearest neighbor Gaussian processes. *arXiv preprint arXiv:2202.01694*, 2022.
- [U3] Jonathan Wenger, **Geoff Pleiss**, Philipp Hennig, John P. Cunningham, and Jacob R. Gardner. Preconditioning for scalable Gaussian process hyperparameter optimization. *arXiv preprint arXiv:2107.00243*, 2022.
- [U4] Jordan Venderley, Michael Matty, Krishnanand Mallayya, Matthew Krogstad, Jacob Ruff, **Geoff Pleiss**, Varsha Kishore, David Mandrus, Daniel Phelan, Lekhanath Poudel, and others. Harnessing interpretable and unsupervised machine learning to address big data from modern x-ray diffraction. *arXiv preprint arXiv:2008.03275*, 2021.

Technical Reports and Workshop Proceedings

- [T1] Martin Jankowiak and **Geoff Pleiss**. Scalable cross validation losses for Gaussian process models. *arXiv preprint arXiv:2105.11535*, 2021.
- [T2] Elliott Gordon-Rodriguez, Gabriel Loaiza-Ganem, **Geoff Pleiss**, and John P. Cunningham. Uses and abuses of the cross-entropy loss: Case studies in modern deep learning. In *NeurIPS “I Can’t Believe It’s Not Better!” Workshop*, 2020.
- [T3] **Geoff Pleiss**, Danlu Chen, Gao Huang, Tongcheng Li, Laurens van der Maaten, and Kilian Q. Weinberger. Memory-efficient implementation of DenseNets. *arXiv preprint arXiv:1707.06990*, 2017.

SELECTED OPEN SOURCE

- 2018- GPyTorch
 <https://gpytorch.ai>
- 2017 Memory Efficient DenseNets
 https://github.com/gpleiss/efficient_densenet_pytorch