

GEOFF PLEISS, CURRICULUM VITAE

Last updated: 6th August, 2024

University of British Columbia
Department of Statistics
2207 Main Mall
Vancouver, BC Canada V6T 1Z4

Homepage: <http://www.geoffpleiss.com>
Email: geoff.pleiss@stat.ubc.ca
Phone: (604) 827-3091
Citizenship: Canada, USA

ACADEMIC POSITIONS AND EDUCATION

- 2023– UNIVERSITY OF BRITISH COLUMBIA (Vancouver, BC, Canada)
Assistant Professor, Department of Statistics (2023–)
Associate Member, Department of Computer Science (2023–)
Centre for Artificial Intelligence Decision-Making and Action (CAIDA)
Artificial Intelligence Methods for Scientific Impact (AIM-SI) Cluster
- 2023– VECTOR INSTITUTE (Toronto, ON, Canada)
CIFAR AI Chair (2024–)
Faculty Member (2023–)
- 2020–2023 COLUMBIA UNIVERSITY (New York, NY, USA)
Postdoctoral Research Scientist, Zuckerman Institute
Supervisor: John P. Cunningham
- 2015–2020 CORNELL UNIVERSITY (Ithaca, NY, USA)
Ph.D., Computer Science (2020)
M.S., Computer Science (2018)
Committee: Kilian Q. Weinberger (chair), Andrew Gordon Wilson, Karthik Sridharan
Dissertation: A Scalable and Flexible Framework for Gaussian Processes via Matrix-Vector Multiplication
- 2009–2013 OLIN COLLEGE OF ENGINEERING (Needham, MA, USA)
B.Sc., Engineering (2013)
Concentration: Computing with Applied Mathematics

OTHER RELEVANT EXPERIENCE

- 2019–2020 ASAPP, INC. (Ithaca, NY, USA)
Research Intern
- 2018 MICROSOFT, INC. (Redmond, WA, USA)
Research Intern
- 2013–2015 PIVOTAL INC. (New York, NY, USA)
Software Engineer

HONOURS AND AWARDS

2024	Canada CIFAR AI Chair
2023	AISTATS Top Reviewer (top 10%)
2022	NeurIPS “I Can’t Believe It’s Not Better” Workshop – Most Surprising Result Award
2022	AISTATS Top Reviewer (top 10%)
2020	NeurIPS Top Reviewer (top 10%)
2019	NeurIPS Top Reviewer (top 50%)
2017	National Science Foundation Graduate Research Fellowship (honorable mention)
2016	National Science Foundation Graduate Research Fellowship (honorable mention)
2012	Barry M. Goldwater Scholarship (honorable mention)
2009–2013	Olin Merit Scholarship (full-tuition recipient)

PUBLICATIONS

** denotes equal author contribution (shared first-authorship).*

Citation Statistics

All statistics are based on Google Scholar, with manual corrections for errors.

Total citations of all publications: 11000+

Total citations of top-three most cited publications: 7800+

Publications (including technical reports) with 100+ citations: 11

Publications (including technical reports) with 10+ citations: 25

Preprints Under Submission

- [U1] Mahdi Ebrahimi Kahou, James Yu, Jesse Perla, and **Geoff Pleiss**. How inductive bias in machine learning aligns with optimality in economic dynamics. *arXiv preprint arXiv:2406.01898*, 2024.
- [U2] Jonathan Wenger, Kaiwen Wu, Philipp Hennig, Jacob R. Gardner, **Geoff Pleiss**, and John P. Cunningham. Computation-aware Gaussian processes: Model selection and linear-time inference. 2024.
- [U3] Natalie Maus, Kyurae Kim, **Geoff Pleiss**, David Eriksson, John P. Cunningham, and Jacob R. Gardner. Approximation-aware Bayesian optimization. *arXiv preprint arXiv:2406.04308*, 2024.
- [U4] Alexandre Bouchard-Côté, Trevor Campbell, **Geoff Pleiss**, and Nikola Surjanovic. MCMC-driven learning. *arXiv preprint arXiv:2402.09598*, 2024.

Refereed Conference Publications

In machine learning, conferences are considered prestigious venues for publication. All venues listed here are highly selective (acceptance rate 20 – 30%) and have peer-reviewing and refereeing processes similar to journals.

- [C1] Agustinus Kristiadi, Felix Strieth-Kalthoff, Marta Skreta, Pascal Poupart, Alán Aspuru-Guzik, and **Geoff Pleiss**. A sober look at LLMs for material discovery: Are they actually good for Bayesian optimization over molecules? In *International Conference on Machine Learning*, 2024.
- [C2] Jinsoo Yoo, Yunpeng Liu, Frank Wood, and **Geoff Pleiss**. Layerwise proximal replay: A proximal point method for online continual learning. In *International Conference on Machine Learning*, 2024.
- [C3] Kaiwen Wu, Jonathan Wenger, Hadyn Jones, **Geoff Pleiss**, and Jacob R. Gardner. Large-scale Gaussian processes via alternating projection. In *Artificial Intelligence and Statistics*, 2024.
- [C4] Andres Potapczynski*, Marc Anton Finzi*, **Geoff Pleiss**, and Andrew Gordon Wilson. CoLA: Exploiting compositional structure for automatic and efficient numerical linear algebra. In *Neural Information Processing Systems*, 2023.
- [C5] Alexandre Capone, Sandra Hirche, and **Geoff Pleiss**. Sharp calibrated Gaussian processes. In *Neural Information Processing Systems*, 2023.
- [C6] Jonathan Wenger, **Geoff Pleiss**, Marvin Pförtner, Philipp Hennig, and John P. Cunningham. Posterior and computational uncertainty in Gaussian processes. In *Neural Information Processing Systems*, 2022.
- [C7] Taiga Abe*, E. Kelly Buchanan*, **Geoff Pleiss**, Richard Zemel, and John P. Cunningham. Deep ensembles work, but are they necessary? In *Neural Information Processing Systems*, 2022.
- [C8] Luhuan Wu, **Geoff Pleiss**, and John P. Cunningham. Variational nearest neighbor Gaussian processes. In *International Conference on Machine Learning*, 2022.
- [C9] Jonathan Wenger, **Geoff Pleiss**, Philipp Hennig, John P. Cunningham, and Jacob R. Gardner. Pre-conditioning for scalable Gaussian process hyperparameter optimization. In *International Conference on Machine Learning*, 2022. [LONG ORAL PRESENTATION—TOP 2% OF SUBMISSIONS].
- [C10] **Geoff Pleiss** and John P. Cunningham. The limitations of large width in neural networks: A deep Gaussian process perspective. In *Neural Information Processing Systems*, 2021.
- [C11] Anthony L. Caterini*, Gabriel Loaiza-Ganem*, **Geoff Pleiss**, and John P. Cunningham. Rectangular flows for manifold learning. In *Neural Information Processing Systems*, 2021.
- [C12] Andres Potapczynski*, Luhuan Wu*, Dan Biderman*, **Geoff Pleiss**, and John P. Cunningham. Bias-free scalable Gaussian processes via randomized truncations. In *International Conference on Machine Learning*, 2021.
- [C13] Luhuan Wu*, Andrew Miller*, Lauren Anderson, **Geoff Pleiss**, David Blei, and John P. Cunningham. Hierarchical inducing point Gaussian process for inter-domain observations. In *Artificial Intelligence and Statistics*, 2021.
- [C14] **Geoff Pleiss**, Martin Jankowiak, David Eriksson, Anil Damle, and Jacob R. Gardner. Fast matrix square roots with applications to Gaussian processes and Bayesian optimization. In *Neural Information Processing Systems*, 2020.
- [C15] **Geoff Pleiss**, Tianyi Zhang, Ethan Elenberg, and Kilian Q. Weinberger. Identifying mislabeled data using the area under the margin ranking. In *Neural Information Processing Systems*, 2020.
- [C16] Martin Jankowiak, **Geoff Pleiss**, and Jacob R. Gardner. Deep sigma point processes. In *Uncertainty in Artificial Intelligence*, 2020.

- [C17] Martin Jankowiak, **Geoff Pleiss**, and Jacob R. Gardner. Parametric Gaussian process regressors. In *International Conference on Machine Learning*, 2020.
- [C18] Yurong You*, Yan Wang*, Wei-Lun Chao*, Divyansh Garg, **Geoff Pleiss**, Bharath Hariharan, Mark Campbell, and Kilian Q. Weinberger. Pseudo-lidar++: Accurate depth for 3d object detection in autonomous driving. In *International Conference on Learning Representations*, 2020.
- [C19] Ke Wang*, **Geoff Pleiss***, Jacob R. Gardner, Stephen Tyree, Kilian Q. Weinberger, and Andrew Gordon Wilson. Exact Gaussian processes on a million data points. In *Neural Information Processing Systems*, 2019.
- [C20] Jacob R. Gardner*, **Geoff Pleiss***, David Bindel, Kilian Q. Weinberger, and Andrew Gordon Wilson. GPyTorch: Blackbox matrix-matrix Gaussian process inference with GPU acceleration. In *Neural Information Processing Systems*, 2018. [SPOTLIGHT PRESENTATION—TOP 4% OF SUBMISSIONS].
- [C21] **Geoff Pleiss**, Jacob R. Gardner, Andrew Gordon Wilson, and Kilian Q. Weinberger. Constant time predictive distributions for Gaussian processes. In *International Conference on Machine Learning*, 2018.
- [C22] Jacob R. Gardner, **Geoff Pleiss**, Ruihan Wu, Andrew Gordon Wilson, and Kilian Q. Weinberger. Product kernel interpolation for scalable Gaussian processes. In *Artificial Intelligence and Statistics*, 2018.
- [C23] **Geoff Pleiss***, Manish Raghavan*, Felix Wu, Jon Kleinberg, and Kilian Q. Weinberger. On fairness and calibration. In *Neural Information Processing Systems*, 2017.
- [C24] Chuan Guo*, **Geoff Pleiss***, Yu Sun*, and Kilian Q. Weinberg. On calibration of modern neural networks. In *International Conference on Machine Learning*, 2017.
- [C25] Paul Upchurch*, Jacob Gardner*, **Geoff Pleiss**, Kavita Bala, Robert Pless, Noah Snavely, and Kilian Q. Weinberger. Deep feature interpolation for image content changes. In *Computer Vision and Pattern Recognition*, 2017.
- [C26] Gao Huang*, Yixuan Li*, **Geoff Pleiss**, Zhuang Liu, John E. Hopcroft, and Kilian Q. Weinberger. Snapshot ensembles: Train 1, get M for free. In *International Conference on Learning Representations*, 2017.

Journal Publications

- [J1] Taiga Abe, E. Kelly Buchanan, **Geoff Pleiss**, and John P. Cunningham. Pathologies of predictive diversity in deep ensembles. *Transactions on Machine Learning Research*, 2024. [FEATURED PAPER—TOP 2% OF SUBMISSIONS].
- [J2] Jordan Venderley, Michael Matty, Krishnanand Mallayya, Matthew Krogstad, Jacob Ruff, **Geoff Pleiss**, Varsha Kishore, David Mandrus, Daniel Phelan, Lekhanath Poudel, and others. Harnessing interpretable and unsupervised machine learning to address big data from modern x-ray diffraction. *Proceedings of the National Academy of Sciences*, 119(24), 2022.
- [J3] Gao Huang*, Zhuang Liu*, **Geoff Pleiss**, Laurens van der Maaten, and Kilian Q. Weinberger. Convolutional networks with dense connectivity. *Transactions on Pattern Analysis and Machine Intelligence*, 2019.
- [J4] James Knighton, **Geoff Pleiss**, Elizabeth Carter, Steven Lyon, M. Todd Walter, and Scott Steinschneider. Potential predictability of regional precipitation and discharge extremes using synoptic-scale climate information via machine learning: An evaluation for the eastern continental United States. *Journal of Hydrometeorology*, 20(5):883–900, 2019.

Technical Reports and Workshop Proceedings

- [R1] Jason Yoo, Dylan Green, **Geoff Pleiss**, and Frank Wood. Online continual learning of video diffusion models from a single video stream. *arXiv preprint arXiv:2406.04814*, 2024.
- [R2] Agustinus Kristiadi, Felix Strieth-Kalthoff, Sriram Ganapathi Subramanian, Vincent Fortuin, Pascal Poupart, and **Geoff Pleiss**. How useful is intermittent, asynchronous expert feedback for Bayesian optimization? In *Symposium on Advances in Approximate Bayesian Inference*, 2024.
- [R3] E. Kelly Buchanan, **Geoff Pleiss**, Yixin Wang, and John P. Cunningham. The effects of ensembling on long-tailed data. In *NeurIPS “Heavy Tails in ML: Structure, Stability, Dynamics” Workshop*, 2023.
- [R4] Taiga Abe*, E. Kelly Buchanan*, **Geoff Pleiss**, and John P. Cunningham. The best deep ensembles sacrifice predictive diversity. In *NeurIPS “I Can’t Believe It’s Not Better!” Workshop*, 2022. [ORAL PRESENTATION].
- [R5] Martin Jankowiak and **Geoff Pleiss**. Scalable cross validation losses for Gaussian process models. *arXiv preprint arXiv:2105.11535*, 2021.
- [R6] Elliott Gordon-Rodriguez, Gabriel Loaiza-Ganem, **Geoff Pleiss**, and John P. Cunningham. Uses and abuses of the cross-entropy loss: Case studies in modern deep learning. In *NeurIPS “I Can’t Believe It’s Not Better!” Workshop*, 2020. [ORAL PRESENTATION].
- [R7] **Geoff Pleiss***, Danlu Chen*, Gao Huang, Tongcheng Li, Laurens van der Maaten, and Kilian Q. Weinberger. Memory-efficient implementation of DenseNets. *arXiv preprint arXiv:1707.06990*, 2017.

SELECTED OPEN SOURCE

Co-Founder and Maintainer

- | | |
|-------|--|
| 2018– | GPyTorch
https://gpytorch.ai |
| 2022– | LinearOperator
https://linear-operator.readthedocs.io |
| 2023– | CoLA (Compositional Linear Algebra)
https://cola.readthedocs.io/ |

PATENTS

- [P1] Tianyi Zhang, Sam Altschul, Kilian Weinberger, Michael Griffiths, and **Geoff Pleiss**. Trend detection via machine learning models, September 2023. US Patent #11,763,230.

GRANTS

- | | |
|------|------------------------------------|
| 2024 | NSERC Discovery (PI) |
| 2024 | NSERC Early Career Supplement (PI) |

INVITED TALKS

June 2024	Adaptive Experimentation Workshop, Meta Inc. (New York, NY, USA) <i>Combining LLMs and Black-Box Optimization for Materials Discovery</i>
June 2024	Statistical Society of Canada Annual Meeting (St. John's, NL, Canada) <i>Ensembles in the Age of Overparameterization: Promises and Pathologies</i>
May 2024	AI In Medicine Meeting, Karolinska Institute (Stockholm, Sweden) <i>Foundation Models for Science: A Case Study on Combining LLMs and Black-Box Optimization for Materials Discovery</i>
April 2024	Academic Seminar, Two Sigma Investments LP (New York, NY, USA) <i>Towards Reliable Neural Networks: What's Within Grasp and What's Out of Reach?</i>
Dec. 2023	Vector Distinguished Talk Series, Vector Institute (Toronto, ON, Canada) <i>Troubling Trajectories for Uncertainty Quantification and Decision Making with Neural Networks</i>
Spring 2022	Job talk, various universities <i>Bridging The Gap Between Deep Learning and Probabilistic Modeling</i>
Oct. 2021	Artificial Intelligence Seminar, University College London (Virtual) <i>Understanding Neural Networks through Gaussian Processes, and Vice Versa</i>
Dec. 2020	Machine Learning for Nuclear Data Workshop (Virtual) <i>GPyTorch: A Scalable and Flexible Framework for Gaussian Processes via Matrix-Vector Multiplication</i>
May 2020	Columbia University (Virtual) <i>GPyTorch: A Scalable and Flexible Framework for Gaussian Processes via Matrix-Vector Multiplication</i>
Nov. 2019	Computer Science Colloquium, Cornell University (Ithaca, NY, USA) <i>From $N = 1,000$ to $N = 1,000,000$: Scaling Gaussian Process Inference with Matrix Multiplication and GPU Acceleration</i>
May 2019	Symposium on Bayesian Optimization, Uber AI (San Francisco, CA, USA) <i>From $N = 1,000$ to $N = 1,000,000$: Scaling Gaussian Process Inference with Matrix Multiplication and GPU Acceleration</i>

CONTRIBUTED TALKS

Aug. 2024	Joint Statistical Meeting (Portland, OR, USA) <i>Task-Aware Scalable Gaussian Processes</i>
Feb. 2024	SIAM UQ Conference (Trieste, Italy) <i>Blurring the Distinction Between Data Collection and Computation in Gaussian Processes</i>
Dec. 2018	Neural Information Processing Systems (Montreal, QC, Canada) <i>GPyTorch: Blackbox Matrix-Matrix Gaussian Process Inference with GPU Acceleration</i>
July 2018	International Conference on Machine Learning (Stockholm, Sweden) <i>Constant Time Predictive Distributions for Gaussian Processes</i>
Aug. 2017	International Conference on Machine Learning (Sydney, Australia) <i>On Calibration of Modern Neural Networks</i>

INVITED LECTURES

May 2024 Swedish NDPIA “AI Applications in Infection Biology” Course (Rånäs, Sweden)
Machine Learning Fundamentals I and II

TEACHING

University of British Columbia

Spring 2025 STAT 547U — Topics in Deep Learning Theory
Fall 2024 STAT 406 — Methods for Statistical Learning
Fall 2023 STAT 520P — Bayesian Optimization

ADVISING AND SUPERVISION

Supervision

2024– Donney Fan (Ph.D.), Department of Computer Science, UBC (co. Mark Schmidt)
2024–2025 Joey Hotz (MS.C.), Department of Statistics, UBC
2024–2025 Timothy Zhou (Undergrad), UBC
2024 Tommy Xu (Undergrad), UBC
2024 Niclas Dern (Intern), Vector Institute
2024 Felix Fu (Intern), Vector Institute

Research Committee Membership (Excluding Direct Supervision)

2024– Nikola Surajonovic (Ph.D.), Department of Statistics, UBC
2024– Qiaoyue Tang (Ph.D.), Department of Computer Science, UBC

Doctoral Thesis Pre-Examiner

2024 Paul E. Chang, Aalto University

PROFESSIONAL SERVICE

Area Chair

International Conference on Machine Learning (2022–2024)
International Conference on Learning Representations (2024)
International Joint Conference on Artificial Intelligence (2023)
Neural Information Processing Systems (2022–2024)

Organizing Committee Member

NeurIPS Workshop on Gaussian Processes, Spatiotemporal Modeling, and Decision-Making Systems (2022)

Virtual Seminar on Gaussian Processes, Spatiotemporal Modeling, and Decision-Making Systems (2022–2023)

Panelist

Scientific Software Development Panel: Dagstuhl Seminar on Probabilistic Numerical Methods (2021)

Conference Reviewer

AAAI Conference on Artificial Intelligence (2017)

Artificial Intelligence and Statistics (2019–2024)

International Conference on Learning Representations (2022)

International Conference on Machine Learning (2019–2021)

Neural Information Processing Systems (2018–2021)

Uncertainty in Artificial Intelligence (2018)

Journal Reviewer

Bernoulli (2022)

Journal of Machine Learning Research (2019–2022)

SIAM/ASA Journal on Uncertainty Quantification (2024)

Transactions on Machine Learning Research (2022–2023)

Transactions on Pattern Analysis and Machine Intelligence (2020–2021)

Workshop Reviewer

NeurIPS Workshop on Gaussian Processes, Spatiotemporal Modeling, and Decision-Making Systems (2024)

NeurIPS “I Can’t Believe It’s Not Better” Workshop (2023)

NeurIPS “Your Model is Wrong: Robustness and Misspecification in Probabilistic Modeling” Workshop (2021)

Other

NeurIPS—workshop proposal reviewer (2024)

MEDIA APPEARANCES

May 2024 Vector Institute Research Blog: [The known unknowns: Vector researcher Geoff Pleiss digs deep into uncertainty to make ML models more accurate](#)

May 2023 “The Ensembles Podcast”

OUTREACH

Fall 2020	LatinX in AI NeurIPS mentorship program
Spring 2018	Cornell "Expand Your Horizons" (STEM workshop for middle school girls)
Spring 2017	Cornell "GRASSHOPR" (After-school CS class at local middle school)
Spring 2016	Cornell "Expand Your Horizons"
Spring 2016	"Code4Kids" (After-school CS class at local elementary school)